

Does a Notch Make a Non-Linearity?

An Empirical Evaluation of HeLU against ReLU and GELU

HeLU activation study

`helu-activation-research/` in `lukeabraham24777/unblock-me-game`

September 2026

Abstract

We evaluate a proposed activation function, HeLU, defined as $\text{HeLU}(x) = x$ everywhere except on the interval $0.5 \leq x \leq 0.6$, where $\text{HeLU}(x) = 0.9x$. The hypothesis under test is that networks trained with HeLU outperform networks trained with ReLU or GELU. We compare the three activations, plus an identity (purely linear) control, under identical architectures, optimisers, and seeds on one-dimensional function regression, two-dimensional non-linearly separable classification, MNIST and Fashion-MNIST with multilayer perceptrons, Fashion-MNIST with a small convolutional network, a depth sweep, and a mechanistic analysis of the trained networks. **The hypothesis is not supported.** On every task whose solution requires a non-linear decision function, HeLU networks perform substantially and statistically significantly worse than ReLU and GELU networks, and are statistically indistinguishable from the linear control. A least-squares linear fit explains a fraction 1.00000 of the variance of a trained HeLU network’s outputs (versus 0.957 for ReLU), and only 2.81% of hidden pre-activations fall inside the notch after training. HeLU is therefore best understood as an approximately linear map with a small, rarely visited perturbation, and it inherits the expressivity limits of deep linear networks. The one setting in which HeLU comes within a point of ReLU and GELU is the convolutional network, where max-pooling supplies the non-linearity that the activation does not, and the linear control comes equally close there. An ablation over sixteen variants of the notch (moved, deepened, or widened) shows that the deficit is a property of the notch design rather than of its particular constants, and a periodic sawtooth function, tested as a further alternative, fails differently again: it is expressive but its constant slope makes the gradient blind to which period a unit is in, and training never leaves chance. All code, raw results, and figures are released alongside this document.

1 Introduction

The choice of activation function sets the family of functions a neural network can represent and shapes the optimisation landscape. ReLU [6] and GELU [2] dominate current practice; the former is cheap and sparse, the latter is smooth and has become the default in transformer architectures. Proposals for new activations are common, and the standard of evidence is empirical: a new function is judged by whether models trained with it reach better accuracy or lower loss than the same models trained with the incumbents [8].

This report applies that standard to a proposed activation, which we refer to as HeLU. It is defined as

$$\text{HeLU}(x) = \begin{cases} 0.9x & \text{if } 0.5 \leq x \leq 0.6, \\ x & \text{otherwise.} \end{cases} \quad (1)$$




Figure 1: Left: ReLU, GELU, and HeLU on $[-2, 2]$. Centre: HeLU near the notch; filled and open circles mark the value taken and not taken at each edge. Right: derivatives. HeLU’s derivative is 1 except on $[0.5, 0.6]$, where it is 0.9.

The claim we set out to verify is that HeLU is *superior* to ReLU and GELU, i.e. that networks trained with HeLU achieve better held-out performance under matched conditions. We fixed the following reading of the outcome before running the full experimental sweep: the hypothesis is supported if HeLU attains higher mean test accuracy (or lower test error) than both ReLU and GELU on the majority of tasks, with the difference significant at $p < 0.05$ under Welch’s t -test [10] across seeds.

Contributions.

1. A controlled comparison of HeLU, ReLU, and GELU, plus an identity control, over eight experiments spanning regression, classification, image benchmarks, and depth (Section 3).
2. A negative result: HeLU is statistically indistinguishable from a linear network on every task we tried and significantly worse than ReLU and GELU whenever the task is non-linear (Section 4).
3. A mechanistic explanation. The notch is visited by roughly 2.81% of pre-activations, and a linear model fits the trained HeLU network almost perfectly (Section 4.5).

2 The HeLU function

Figure 1 plots the three activations and their derivatives. Several properties of (1) follow directly from the definition.

- **It is almost the identity.** $|\text{HeLU}(x) - x| = 0.1|x| \leq 0.06$, and this deviation is non-zero only on an interval of width 0.1.
- **It is discontinuous.** At $x = 0.5$ the function jumps from 0.5 to 0.45; at $x = 0.6$ it jumps from 0.54 to 0.6. Neither ReLU nor GELU has a discontinuity.
- **Its derivative is nearly constant.** $\text{HeLU}'(x) = 1$ outside the notch and 0.9 inside; automatic differentiation reproduces exactly this (Figure 1, right). Gradients therefore flow through HeLU layers essentially unchanged, as in a linear network.
- **It is non-polynomial.** By the theorem of Leshno et al. [5], a single-hidden-layer network with HeLU activations is a universal approximator in the limit of infinite width. The theorem, however, says nothing about how many units, how large the weights, or how difficult the optimisation problem must be. For HeLU the non-linear budget of each unit is a perturbation of amplitude at most 0.06 confined to a band of width 0.1, so expressing a non-linear function

of order-one amplitude requires either very many units whose pre-activations land in the notch or very large output weights.

The natural null model for HeLU is therefore a *deep linear network*, whose behaviour is well understood [9]: regardless of depth or width it can only represent a single affine map of its input. We include the identity activation as an explicit control so that this null can be tested directly rather than assumed.

3 Experimental setup

All experiments use PyTorch [7] on CPU with the Adam optimiser [3]. Within each experiment the architecture, learning rate, batch size, number of steps or epochs, data ordering, and parameter-initialisation seed are identical across activations; the activation module is the only thing that changes. Hyper-parameters were fixed in advance to common defaults and were not tuned per activation. Results are reported as mean $\pm$ sample standard deviation across seeds, and differences involving HeLU are tested with Welch’s unequal-variance *t*-test.

E1: Function shape. Analytic evaluation of f and f' (Figure 1).

E2: 1-D regression. An MLP with two hidden layers of 64 units maps $x \in [-2, 2]$ to one of four targets: $\sin 3x$, $|x|$, x^2 , and the step $\mathbf{1}[x > 0]$. Training uses 512 points with Gaussian noise ($\sigma = 0.05$), 3000 full-batch Adam steps at learning rate 3×10^{-3} ; the test MSE is computed on 2000 noise-free points. 5 seeds.

E3: 2-D classification. Two-moons and two-spirals datasets (600 training points, 4000 test points), same MLP with a two-way softmax output, 3000 steps. 5 seeds.

E4: MNIST and Fashion-MNIST, MLP. 784–256–256–10 MLP, batch size 128, learning rate 10^{-3} , 5 epochs. 5 seeds each [4, 11].

E5: Fashion-MNIST, CNN. Two 3×3 convolutions (32 and 64 channels), each followed by the activation and 2×2 max-pooling, then a 128-unit fully connected layer with the activation and a linear classifier. Same optimiser settings, 5 epochs, 3 seeds.

E6: Depth sweep. Fashion-MNIST MLP with 128-unit hidden layers and 1, 2, 4, or 8 hidden layers, 3 epochs, 3 seeds.

E7: Linearity and notch occupancy. For the Fashion-MNIST MLP we measure (i) the R^2 of the best least-squares affine map from the 784-dimensional input to the network’s 10 logits, computed on 2000 test images, and (ii) the fraction of hidden pre-activations lying in $[0.5, 0.6]$, both at initialisation and after 3 epochs of training. 3 seeds.

E8: Kernel cost. Wall time of forward plus backward through the activation alone on a 4×10^6 -element tensor, 20 repetitions.

E9: Notch variants. Sixteen variants of (1) obtained by moving the notch ($[0.5, 0.6]$, $[0.1, 0.2]$, $[-0.05, 0.05]$), deepening it (scale 0.9, 0.8, 0.5, 0.2), widening it ($[0.5, 1.0]$, $[0.5, 2.0]$, $[0.5, \infty)$ at scale 0.2), or replacing it with a wide symmetric band around zero ($[-0.68, 0.68]$ at scale 0.5). Each variant is run under the E2 protocol on $\sin 3x$ and x^2 , the E3 protocol on moons and spirals, and the E4 Fashion-MNIST MLP protocol for 3 epochs, all with 3 seeds.

E10: Sawtooth function. The function that, on each interval $[k, k + 2)$ for $k \in 2\mathbb{Z}$, rises linearly from $(k, 0)$ to $(k + 2, 1)$, i.e. $T(x) = x/2 - \lfloor x/2 \rfloor$. It is run under the E9 protocol, and we additionally record the R^2 of a linear fit to the trained Fashion-MNIST network, a learning-rate sweep over $\{10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}\}$ with the standard deviation of first-layer pre-activations before and after training, and a descent-direction test: on 50 mini-batches at initialisation, a plain gradient step of size η is taken and we record whether the batch loss actually decreased.




Figure 2: 1-D regression (E2). Top: the seed-0 fit of each network against the target (dotted). Bottom: test MSE on a log scale, bars are means over 5 seeds, dots are individual seeds. HeLU and the linear control coincide almost exactly.

Table 1: 1-D regression test MSE (mean $\pm$ sd over 5 seeds). Lower is better.

Target	ReLU	GELU	HeLU	Linear
$\sin(3x)$	0.0003 ± 0.0002	0.0002 ± 0.0003	0.3695 ± 0.0518	0.4364 ± 0.0093
$ x $	0.0001 ± 0.0000	0.0001 ± 0.0000	0.2792 ± 0.0193	0.3409 ± 0.0071
x^2	0.0001 ± 0.0000	0.0000 ± 0.0000	1.1113 ± 0.0634	1.4361 ± 0.0297
$\mathbf{1}[x > 0]$	0.0012 ± 0.0005	0.0021 ± 0.0002	0.0443 ± 0.0092	0.0646 ± 0.0012

4 Results

4.1 Function approximation

Figure 2 and Table 1 summarise E2. ReLU and GELU networks fit all four targets closely, with test MSE of order 10^{-4} to 10^{-3} . The HeLU network does *not* collapse to the linear control’s straight line: with full-batch training on a one-dimensional input, gradient descent does find the notch and uses it. What it builds with it, however, is a jagged staircase. Because the only non-linear features of HeLU are its two jump discontinuities, the network can represent vertical jumps but not smooth curvature, and the fit consists of a linear trend decorated with irregular steps of amplitude comparable to the notch itself. The resulting test MSE is three orders of magnitude worse than ReLU or GELU on $\sin 3x$, $|x|$, and x^2 (Table 1), and only 15–25% better than the linear control. The step target is the most favourable case for HeLU, since a discontinuity is exactly what it can express, and even there its error is 30–40 $\times$ that of ReLU and GELU.




Figure 3: Decision functions on two-moons (top) and two-spirals (bottom), E3. Colour is the predicted probability of class 1; the black curve is the 0.5 level set. Titles give mean test accuracy over 5 seeds. HeLU produces the linear control’s straight boundary with small, spiky perturbations where hidden pre-activations cross the notch.

Table 2: 2-D classification test accuracy (mean $\pm$ sd over 5 seeds) and Welch p -values for the difference between HeLU and each non-linear baseline.

Dataset	Test accuracy				p (HeLU vs.)	
	ReLU	GELU	HeLU	Linear	ReLU	GELU
moons	0.9873 ± 0.0017	0.9845 ± 0.0016	0.8803 ± 0.0078	0.8769 ± 0.0032	$< 10^{-4}$	$< 10^{-4}$
spirals	0.9962 ± 0.0016	0.9963 ± 0.0024	0.6303 ± 0.0121	0.6486 ± 0.0061	$< 10^{-4}$	$< 10^{-4}$

4.2 Non-linearly separable classification

Two-moons and two-spirals are the canonical tests of whether a model can learn a curved decision boundary. Figure 3 shows that ReLU and GELU do, reaching 98.7% and 98.4% on moons and 99.6% and 99.6% on spirals. HeLU reaches 88.0% on moons and 63.0% on spirals, against 87.7% and 64.9% for the linear control (Table 2). Its decision boundary is the linear control’s straight line with small, spiky perturbations where a hidden unit’s pre-activation crosses the notch, visible as faint striations in Figure 3. On moons these perturbations buy a fraction of a percentage point; on spirals they cost one, and HeLU is slightly *worse* than a linear classifier. In neither case does HeLU approach the non-linear baselines, and both differences from ReLU and GELU are significant at $p < 10^{-4}$.

4.3 Image classification

Table 3 and Figures 4 and 5 report E4 and E5. On the MNIST MLP, ReLU and GELU reach 97.57% and 97.91% while HeLU reaches 91.50%, essentially the accuracy of a multinomial logistic regression, which is what a deep linear network trained with cross-entropy computes (linear control:




Figure 4: Training dynamics on the image benchmarks (E4, E5). Top: training loss, averaged over seeds after a 25-step moving average. Bottom: test accuracy after each epoch, shaded bands are ± 1 sd over seeds.

Table 3: Image benchmarks: final test accuracy (mean $\pm$ sd; 5 seeds for the MLPs, 3 for the CNN) and Welch p -values for HeLU against each other activation.

Benchmark	Test accuracy (mean $\pm$ sd)				p (HeLU vs.)		
	ReLU	GELU	HeLU	Linear	ReLU	GELU	Linear
MNIST, MLP	0.9757 $\pm$ 0.0021	0.9791 $\pm$ 0.0011	0.9150 $\pm$ 0.0035	0.9150 $\pm$ 0.0036	$< 10^{-4}$	$< 10^{-4}$	0.98
Fashion-MNIST, MLP	0.8814 $\pm$ 0.0037	0.8816 $\pm$ 0.0036	0.8333 $\pm$ 0.0029	0.8333 $\pm$ 0.0029	$< 10^{-4}$	$< 10^{-4}$	1
Fashion-MNIST, CNN	0.9148 $\pm$ 0.0025	0.9196 $\pm$ 0.0005	0.9035 $\pm$ 0.0050	0.9032 $\pm$ 0.0041	0.0395	0.0293	0.933

91.50%). On the Fashion-MNIST MLP the gap is similar: 88.14% and 88.16% against 83.33% (linear: 83.33%). In both cases the HeLU-versus-ReLU and HeLU-versus-GELU differences are significant at the $p < 0.05$ level, whereas HeLU versus the linear control is not. The training-loss curves in Figure 4 show that HeLU does not optimise faster in early training either: its loss tracks the linear control’s from the first steps and plateaus at the same, higher, level.

The convolutional network narrows the gap without closing it. Here HeLU reaches 90.35%, about one point below ReLU (91.48%, $p = 0.04$) and GELU (91.96%, $p = 0.03$). But the linear control also reaches 90.32%, and HeLU is again indistinguishable from it ($p = 0.93$). The reason the gap is small is that the architecture contains a second non-linearity, the 2×2 max-pooling after each convolution. Max-pooling alone is enough to make the convolutional network a competent non-linear classifier on this dataset, so the activation function’s contribution is largely masked. The CNN result is therefore mostly evidence about max-pooling, not about HeLU; a network with no activation at all obtains the same accuracy.




Figure 5: Final test accuracy on the image benchmarks. Bars are means, whiskers are ± 1 sd, dots are seeds. Note the truncated y -axes.




Figure 6: Fashion-MNIST MLP test accuracy versus number of hidden layers (E6), 3 epochs, 3 seeds.

4.4 Depth

If the notch were being used as a non-linearity, stacking more HeLU layers would compound it and accuracy would change with depth in a way that differs from the linear control. Figure 6 and Table 4 show no such divergence: HeLU’s accuracy is flat from one to four hidden layers and degrades at eight, and it coincides with the linear control to within 0.05 percentage points at every depth. (ReLU and GELU also degrade slightly at depth 8 under this short, untuned 3-epoch budget, but they remain 3–4 points above the linear pair throughout.) A product of affine maps is an affine map, and adding HeLU layers does not change that in practice.

4.5 Mechanism: HeLU networks are linear

Two direct measurements explain the results above (Figure 7, Table 5). First, an affine function of the raw pixels explains an R^2 of 1.00000 of the variance in a trained HeLU network’s logits, against 0.957 for ReLU. The HeLU network is linear to within measurement precision. Second, only about 2.81% of hidden pre-activations lie inside the notch after training, and training does not increase this fraction relative to initialisation. Gradient descent has no incentive to move pre-activations into the notch: doing so changes the unit’s output by at most 0.06 and its gradient by 10%, which is too small a lever to be worth the loss incurred on the way. The notch is not so much unhelpful as unused.

Table 4: Depth sweep on Fashion-MNIST (mean $\pm$ sd over 3 seeds).

Hidden layers	ReLU	GELU	HeLU	Linear
1	0.8667 ± 0.0033	0.8703 ± 0.0027	0.8340 ± 0.0017	0.8341 ± 0.0020
2	0.8700 ± 0.0029	0.8703 ± 0.0017	0.8340 ± 0.0025	0.8341 ± 0.0024
4	0.8621 ± 0.0089	0.8653 ± 0.0060	0.8325 ± 0.0023	0.8321 ± 0.0019
8	0.8548 ± 0.0038	0.8617 ± 0.0048	0.8211 ± 0.0062	0.8212 ± 0.0066




Figure 7: E7. Left: $1 - R^2$ of the best affine fit to each trained network’s logits (log scale; lower means more linear). Right: percentage of HeLU pre-activations that fall inside the notch, per hidden layer, before and after training.

4.6 Computational cost

Table 6 shows that HeLU is not cheaper either. Implemented as a masked `where` with two comparisons, its forward-plus-backward kernel is about $6\times$ slower than ReLU’s and $3\times$ slower than GELU’s in our PyTorch CPU measurement. End-to-end epoch time is dominated by the linear layers, so the difference is small in practice, but there is no efficiency argument for HeLU.

4.7 Variants of the notch

A natural objection to the results so far is that the constants in (1) are arbitrary and a different notch might work. E9 tests this directly (Figure 8, Table 7). Four findings stand out.

Moving the notch does nothing. Placing the notch at $[0.1, 0.2]$ or straddling zero at $[-0.05, 0.05]$ gives, for every scale, results indistinguishable from the linear control on all five tasks. The original position at $[0.5, 0.6]$ is, if anything, the best of the three narrow positions, because a small fraction of pre-activations naturally land there at initialisation.

Deepening the notch helps only on toy problems. Lowering the scale from 0.9 to 0.2 at $[0.5, 0.6]$ reduces the $\sin 3x$ MSE from 0.34 to 0.12 and raises spirals accuracy from 64% to 77%, but ReLU and GELU are at 3×10^{-4} and 99.6% on the same tasks. On Fashion-MNIST none of the twelve narrow-notch variants moves more than 0.1 percentage points from the linear control’s 83.3%. A deeper notch makes each jump larger but does not make the network more inclined to use it.

Widening the notch changes the character of the function. The three variants that widen the band at scale 0.2 behave differently from all the others. On the toy problems they close most of

Table 5: Linearity of trained Fashion-MNIST MLPs and notch occupancy (means over 3 seeds). The notch-occupancy column for ReLU, GELU, and Identity reports the fraction of *their* pre-activations that happen to fall in $[0.5, 0.6]$ and is included for reference only.

Activation	Test acc.	R^2 of linear fit		% pre-act. in $[0.5, 0.6]$	
		init	trained	init	trained
ReLU	0.8707	0.92093	0.95737	2.65	1.26
GELU	0.8740	0.95841	0.95385	2.41	1.35
HeLU	0.8328	0.99918	1.00000	3.44	2.81
Identity (linear control)	0.8328	1.00000	1.00000	3.45	2.80

Table 6: Cost. Rows 1–3: median wall-clock seconds per training epoch over seeds (4 CPU threads on a shared machine; the median is used because a few runs overlapped with other jobs). Last row: forward plus backward time of the activation kernel alone, measured on an otherwise idle machine.

	ReLU	GELU	HeLU	Linear
MNIST, MLP	1.5	1.6	1.8	1.5
Fashion-MNIST, MLP	1.6	1.6	1.7	1.4
Fashion-MNIST, CNN	12.6	13.5	17.2	11.7
Activation kernel, fwd+bwd, 4M elements (ms)	1.1	2.3	6.8	0.5

the gap to ReLU: $[0.5, 1.0] \times 0.2$ reaches 98.4% on moons and 97.9% on spirals and fits $\sin 3x$ with MSE 0.004. On Fashion-MNIST, however, the bounded bands still fail: $[0.5, 1.0]$ is at the linear level and $[0.5, 2.0]$ gains only half a point. Only the unbounded variant $[0.5, \infty) \times 0.2$, which is no longer a notch at all but a two-slope piecewise-linear function with a downward jump at 0.5, reaches 86.36%; that is a real non-linearity, though still 1.8 points below ReLU. Table 8 shows why. After training, only 13.5% of the $[0.5, 1.0]$ variant’s pre-activations sit inside its band and the network is still linear to $R^2 = 0.9992$; the $[0.5, 2.0]$ band is occupied by 31% and R^2 drops to 0.992; the half-line $[0.5, \infty)$ is occupied by 55% of units and R^2 falls to 0.973, comparable to the 0.957 of a ReLU network. A bounded band, however deep, is an “identity plus local dent” whose effect averages out over a unit’s input distribution unless the network concentrates that distribution inside the band, and gradient descent on a high-dimensional problem does not do so. A function that behaves differently on an entire half-line is what forces every unit to be non-linear, and that is exactly the property that ReLU and GELU have and HeLU lacks.

A wide band around zero is escaped rather than used. The symmetric variant $[-0.68, 0.68] \times 0.5$ is the strongest test of the “too small, too rarely visited” explanation, because it is neither: at initialisation 78–98% of Fashion-MNIST pre-activations lie inside it (Table 8). On the toy problems it behaves like the widened notches (98.7% on moons, 97.3% on spirals, $\sin 3x$ MSE 0.005). On Fashion-MNIST it reaches 83.0%, again the linear level, and after training the band occupancy has fallen to about 40% while the network remains linear to $R^2 = 0.9997$. Both halves of this function are linear ($0.5x$ inside, x outside), so a unit is non-linear only if its input distribution straddles ± 0.68 ; the measurements are consistent with training sorting units into one regime or the other rather than exploiting the edges. Where a bounded band ends, the function returns to the identity, and the network can always find a linear solution by arranging for that.




Figure 8: E9: sixteen variants of the HeLU notch, written as $[lo, hi] \times scale$. The first row is the original HeLU. MSE panels (log scale): bars are means over 3 seeds. Accuracy panels (truncated axes): the large marker is the mean and the short line spans the seeds. Small black dots are individual seeds, and the vertical lines are the ReLU, GELU, and linear-control means from the main experiments. Colour groups variants by notch position; yellow marks the widened notches and violet the symmetric band $[-0.68, 0.68] \times 0.5$.

4.8 A sawtooth function

The final function we consider is periodic: on each interval $[k, k + 2)$ with k an even integer, the output is a straight line from $(k, 0)$ to $(k + 2, 1)$, i.e. $T(x) = x/2 - \lfloor x/2 \rfloor$ (Figure 9, left). Its derivative is $1/2$ everywhere except at the even integers, where it drops by 1. Unlike the notch, T is bounded, strongly non-linear, and non-monotone, so a priori it is the kind of function that could give a network real expressive power.

In practice it does not train at all. At the default learning rates, the sawtooth network reaches 50.2% on moons, 54.6% on spirals, and 10.1% on Fashion-MNIST, which is chance on every task, and its $\sin 3x$ MSE (0.49) is worse than the linear control’s (Table 9, Figure 9). The learning-rate sweep in Table 10 rules out a simple step-size problem: on Fashion-MNIST the training loss stays at $\ln 10$ for every learning rate from 10^{-2} to 10^{-5} , and the best result on any task is 71% on moons at $\eta = 10^{-5}$, still far below the linear control.

The failure is not one of gradient *direction* in the local sense. The descent test shows that at initialisation a small gradient step reduces the batch loss on 84–100% of mini-batches, against 100% for ReLU, HeLU, and the identity. The gradient is locally correct; it is globally uninformative. Because T has the same slope $1/2$ at every point of every period, the gradient with respect to

Table 7: E9: HeLU variants (means over 3 seeds; Fashion-MNIST also shows sd). MSE lower is better, accuracy higher is better. Baselines from the main experiments are repeated at the bottom.

Notch	Scale	$\sin 3x$ MSE	x^2 MSE	Moons (%)	Spirals (%)	Fashion-MNIST (%)
[0.5, 0.6]	0.9	0.342	1.075	87.5	63.7	83.29 $\pm$ 0.31
[0.5, 0.6]	0.8	0.310	1.080	87.6	65.9	83.26 $\pm$ 0.30
[0.5, 0.6]	0.5	0.221	0.894	91.2	72.8	83.23 $\pm$ 0.32
[0.5, 0.6]	0.2	0.117	0.826	91.8	76.5	83.28 $\pm$ 0.30
[0.1, 0.2]	0.9	0.433	1.433	87.4	64.3	83.27 $\pm$ 0.31
[0.1, 0.2]	0.8	0.431	1.424	87.4	65.2	83.26 $\pm$ 0.29
[0.1, 0.2]	0.5	0.408	1.405	87.9	71.4	83.26 $\pm$ 0.28
[0.1, 0.2]	0.2	0.382	1.377	88.7	77.7	83.21 $\pm$ 0.25
[-0.05, 0.05]	0.9	0.435	1.449	87.6	64.8	83.26 $\pm$ 0.30
[-0.05, 0.05]	0.8	0.436	1.449	87.6	64.0	83.28 $\pm$ 0.30
[-0.05, 0.05]	0.5	0.431	1.449	87.6	66.3	83.21 $\pm$ 0.27
[-0.05, 0.05]	0.2	0.431	1.445	87.8	67.2	83.24 $\pm$ 0.25
[0.5, 1]	0.2	0.004	0.016	98.4	97.9	83.22 $\pm$ 0.37
[0.5, 2]	0.2	0.008	0.018	98.4	98.0	83.84 $\pm$ 0.30
[0.5, ∞]	0.2	0.005	0.001	98.7	98.3	86.36 $\pm$ 0.33
[-0.68, 0.68]	0.5	0.005	0.042	98.6	97.3	83.01 $\pm$ 0.51
ReLU		0.0003	0.0001	98.7	99.6	88.14
GELU		0.0002	0.0000	98.4	99.6	88.16
Identity (linear control)		0.4364	1.4361	87.7	64.9	83.33

a pre-activation never depends on where in the period that pre-activation lies, so nothing in the gradient signals that a unit has wrapped around a discontinuity. The last column of Table 10 shows the consequence: at $\eta = 10^{-3}$ the standard deviation of first-layer pre-activations grows from 0.6 at initialisation to about 257 after three epochs, spanning hundreds of periods of T , and the drift scales with the learning rate. Each step is a descent step for the current mini-batch, but the weights drift monotonically through period after period, the features seen by the next layer change chaotically from step to step, and no consistent representation is ever formed. The one-dimensional fit in Figure 9 shows the same pathology in miniature: a dense, high-frequency oscillation rather than a curve.

The sawtooth therefore fails for a third, distinct reason. HeLU is too nearly linear for its non-linearity to matter; the widened notches are escaped because their non-linearity is local; and the sawtooth is expressive enough but its periodic structure decouples the gradient from the position of each unit within a period, so gradient descent cannot navigate it. A well-behaved activation needs a derivative that carries information about where the unit is operating, which is exactly what the changing slopes of ReLU and GELU provide.

5 Discussion

Why HeLU fails. A useful activation function must do two things: give the network a non-linear family of functions to search over, and make that family reachable by gradient descent. HeLU does the first only in the technical sense of Leshno et al. [5], and the second not at all. Its total non-linear content, a 10% dip on a band of width 0.1, is an order of magnitude smaller than the scale of typical pre-activations at initialisation [1], and gradient descent finds nothing to gain by

Table 8: Band occupancy and linearity for the original notch, the three widened variants, and the symmetric band (Fashion-MNIST MLP, 3 epochs, means over 3 seeds). “% pre-act. in band” is the fraction of hidden pre-activations inside $[lo, hi]$, averaged over the two hidden layers.

Notch	Scale	Test acc. (%)	% pre-act. in band		
			init	trained	R^2 trained
$[0.5, 0.6]$	0.9	83.28	3.4	2.8	1.0000
$[0.5, 1]$	0.2	83.17	9.2	13.5	0.9992
$[0.5, 2]$	0.2	83.75	10.4	30.7	0.9923
$[0.5, \infty]$	0.2	86.44	10.4	54.7	0.9734
$[-0.68, 0.68]$	0.5	83.01	87.8	39.6	0.9997




Figure 9: E10. From left: the sawtooth $T(x) = x/2 - \lfloor x/2 \rfloor$ (filled and open circles mark the closed and open ends of each ramp); the seed-0 fit to $\sin 3x$ (target dotted); the spirals decision function; test accuracy versus learning rate on the three classification tasks (dotted line is chance on Fashion-MNIST); and the standard deviation of first-layer pre-activations on Fashion-MNIST before and after training (the shaded band is one period of T).

concentrating pre-activations there. The result is a network that is, for all practical purposes, linear, and it performs exactly as a linear network does: adequately on MNIST, where a linear classifier already reaches about 91.5%, and poorly everywhere a curved boundary is required. The one place we saw the notch being used at all, the one-dimensional regression tasks with full-batch training, it produced a jagged staircase of jumps rather than a useful curve.

The CNN caveat. The Fashion-MNIST CNN result deserves emphasis because it is the kind of evidence that could be mistaken for support. HeLU comes within about a point of ReLU on this benchmark only because max-pooling makes the activation function nearly irrelevant, and the identity control demonstrates this by landing at the same accuracy. Any evaluation of a new activation should include a no-activation control for this reason.

What would change the conclusion. Section 4.7 already answers the obvious follow-ups. Moving the notch does nothing; deepening it helps only on one- and two-dimensional toy problems; and the only modification that produces a non-linear network on Fashion-MNIST is to extend the band to the whole half-line $[0.5, \infty)$, at which point the function is a two-slope piecewise linear map with a jump, not a notch, and it still trails ReLU by about two points. We did not tune learning rates per activation; given that HeLU’s training curves coincide with the linear control’s from the first steps, we do not expect tuning to change the ordering.

Table 9: E10: the sawtooth against the main baselines at the default learning rates (means over 3 seeds; Fashion-MNIST also shows sd). R^2 is the linear-fit coefficient of the trained Fashion-MNIST network.

Activation	$\sin 3x$ MSE	x^2 MSE	Moons (%)	Spirals (%)	Fashion-MNIST (%)	R^2
Sawtooth	0.4904	1.4809	50.2	54.6	10.10 ± 0.10	0.3939
ReLU	0.0003	0.0001	98.7	99.6	88.14	0.9574
GELU	0.0002	0.0000	98.4	99.6	88.16	0.9538
HeLU	0.3695	1.1113	88.0	63.0	83.33	1.0000
Identity (linear control)	0.4364	1.4361	87.7	64.9	83.33	1.0000

Table 10: E10 diagnostics. Left: sawtooth learning-rate sweep (3 seeds; Fashion-MNIST after 3 epochs, chance is 10%; “final train loss” is the last 20 mini-batches, chance is $\ln 10 \approx 2.303$). Right: percentage of 50 initial mini-batches on which a gradient step of size η reduced the batch loss.

Learning rate	Moons (%)	Spirals (%)	Fashion-MNIST (%)	Final train loss	Layer-1 pre-act. std	Activation	$\eta = 0.001$	$\eta = 0.01$	$\eta = 0.1$
10^{-2}	50.8	50.4	9.9	2.361	$0.57 \rightarrow 2005.5$	Sawtooth	84	94	100
10^{-3}	51.3	57.4	10.2	2.311	$0.57 \rightarrow 257.3$	ReLU	100	100	100
10^{-4}	67.0	67.4	10.1	2.306	$0.57 \rightarrow 28.0$	HeLU	100	100	100
10^{-5}	71.1	58.3	10.1	2.309	$0.57 \rightarrow 2.9$	Identity	100	100	100

Limitations. All experiments are small-scale and CPU-bound: two image datasets of 28×28 greyscale images, networks of at most a few hundred thousand parameters, at most 5 epochs, and 3 to 5 seeds per configuration. The CIFAR-10 download was blocked in our environment. Larger-scale experiments could in principle behave differently, but the mechanism identified in Section 4.5 does not depend on scale: pre-activations avoid the notch because the gradient signal for entering it is negligible, and that is a property of the function, not of the dataset.

6 Conclusion

We set out to verify that HeLU, the identity function with a $0.9\times$ notch on $[0.5, 0.6]$, is a superior activation to ReLU and GELU. Under matched conditions across regression, two-dimensional classification, MNIST, and Fashion-MNIST, it is not. HeLU networks are statistically indistinguishable from purely linear networks and significantly worse than ReLU and GELU networks on every task that requires a non-linear solution. The mechanism is simple: the notch is too small and too rarely visited to function as a non-linearity, so the network stays linear. The only benchmark on which HeLU comes close to the baselines is one in which max-pooling supplies the non-linearity, and a network with no activation at all comes equally close there. An ablation over sixteen variants confirms that the problem is the notch design rather than its constants: no bounded notch, at any position or depth we tried, produces a non-linear network on Fashion-MNIST, and the one variant that does is no longer a notch. A sawtooth function $x/2 - \lfloor x/2 \rfloor$, tested as a further alternative, stays at chance on every classification task at every learning rate we tried, because its constant slope gives the gradient no information about which period each unit is operating in.

Reproducibility

The code, raw JSON results for every run, figures, and this document are in the repository directory `helu-activation-research/`. Running `run_all.py` regenerates the results (about 40 minutes on

four CPU cores for the main experiments, plus about 30 minutes for the variants), `make_figures.py` regenerates every figure and table in this paper from them, and `make` in the `paper/` directory rebuilds the PDF.

References

- [1] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on ImageNet classification. In *IEEE International Conference on Computer Vision (ICCV)*, 2015.
- [2] Dan Hendrycks and Kevin Gimpel. Gaussian error linear units (GELUs). *arXiv preprint arXiv:1606.08415*, 2016.
- [3] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- [4] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [5] Moshe Leshno, Vladimir Ya. Lin, Allan Pinkus, and Shimon Schocken. Multilayer feedforward networks with a nonpolynomial activation function can approximate any function. *Neural Networks*, 6(6):861–867, 1993.
- [6] Vinod Nair and Geoffrey E. Hinton. Rectified linear units improve restricted Boltzmann machines. In *Proceedings of the 27th International Conference on Machine Learning (ICML)*, 2010.
- [7] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, et al. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [8] Prajit Ramachandran, Barret Zoph, and Quoc V. Le. Searching for activation functions. In *International Conference on Learning Representations (ICLR) Workshop*, 2018.
- [9] Andrew M. Saxe, James L. McClelland, and Surya Ganguli. Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. In *International Conference on Learning Representations (ICLR)*, 2014.
- [10] B. L. Welch. The generalization of ‘student’s’ problem when several different population variances are involved. *Biometrika*, 34(1–2):28–35, 1947.
- [11] Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-MNIST: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.